

Technical Write-Up: Document Intelligence Assistant

1. Overview

The Document Intelligence Assistant is a full-stack GenAI-based application designed to semantically understand documents uploaded by users and answer natural language questions based on their content. It combines document preprocessing, embedding generation, vector-based semantic search, and large language model reasoning in a streamlined Retrieval-Augmented Generation (RAG) pipeline.

2. Embedding Model: HuggingFace **MiniLM-L6-v2**

- We chose **all-MiniLM-L6-v2** from HuggingFace due to its:
 - Lightweight size (~80MB) with excellent performance
 - Compatibility with local CPU/GPU inference
 - Proven performance in semantic search and sentence-level embeddings
 - It enabled fast and meaningful chunk embeddings without requiring external APIs.
-

3. Chunking Strategy

- **Chunker Used:** **RecursiveCharacterTextSplitter**
- **Chunk Size:** 500 characters
- **Overlap:** 50 characters
- **Why Recursive?**
 - Maintains semantic cohesion within text blocks
 - Efficient fallback from paragraphs → sentences → characters

- Prevents mid-sentence breaks and enhances embedding quality
-

4. Prompt Engineering Approach

- The retrieved chunks are inserted into a system-style prompt sent to Groq's LLM.

Prompt Template:

Use the following context to answer the user's question.

Context:

[retrieved chunks go here]

Question: [user input]

Answer:

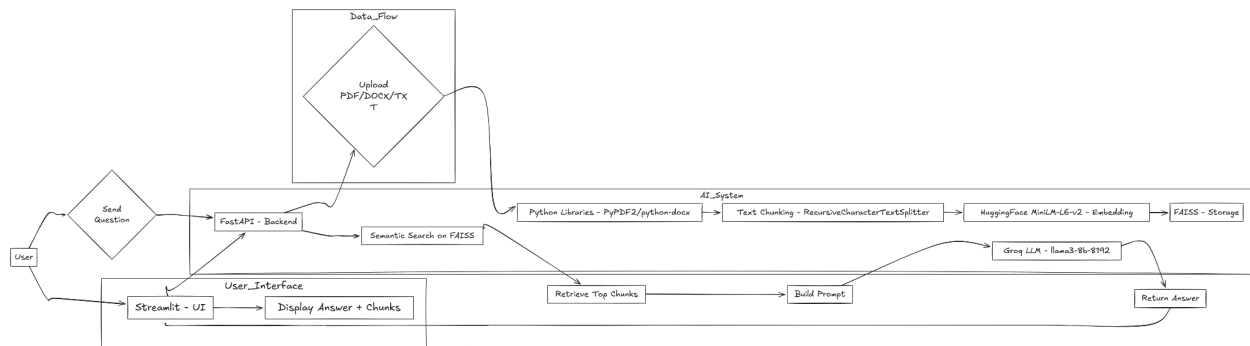
- - This structure ensures:
 - Grounded reasoning
 - No hallucinations
 - Context-sensitive answers
-

5. RAG Architecture Overview

- Users upload `.pdf`, `.docx`, or `.txt` files via Streamlit.
- Files are sent to the FastAPI backend where:
 - Text is extracted (using PyPDF2, python-docx, or file reading)
 - Text is chunked and embedded

- Embeddings are stored in a FAISS vector store
- When a question is asked:
 - FAISS performs similarity search
 - Top relevant chunks are selected
 - Prompt is constructed and sent to Groq's [llama3-8b-8192](#)
 - Answer is returned and shown with context and confidence score

[System Architecture Diagram]



6. Challenges Faced & Solutions

● Chroma Crash on Windows

- **Issue:** Persistent [ConnectionResetError 10054](#) with Chroma
- **Fix:** Switched to FAISS for local vector storage — lightweight and stable

● HuggingFace Embedding Deprecation

- **Issue:** Warning for deprecated embedding class
- **Fix:** Switched to new [langchain-huggingface](#) interface

● FastAPI → Streamlit disconnection

- **Issue:** Streamlit failed silently on file uploads
- **Fix:** Added response delay (`time.sleep(0.1)`) + removed `.persist()` usage

● User confusion on context

- **Issue:** Lack of transparency on how answers are formed
- **Fix:** Exposed retrieved chunks in the frontend with simulated confidence score

● LLM hard fail without fallback

- **Issue:** If Groq API call failed, entire chat broke
- **Fix:** Added fallback mechanism to return a graceful message from a dummy LLM or backup route

7. Outcome

- The app now supports real-time semantic Q&A over uploaded documents.
- Stable, testable, and API-free (except for Groq)
- Compatible across OS platforms due to FAISS choice
- Improved UX with confidence scores and fallback logic

8. Future Enhancements

- Multi-modal support (OCR-based image documents)
- Response streaming from LLM
- Conversation memory
- Document summarization feature

- Query rewriting to expand vague questions
- Multiple LLM fallback (OpenRouter, Gemini, etc.)

Submitted by: Aayush Dharpure

Project: GenAI Document Intelligence Assistant